

# AI for Bharat Hackathon

Powered by 



**Team Name : Assemblers**

**Team Leader Name : Sandeep Udeg**

**Problem Statement : [Professional Track] AI for Rural Innovation & Sustainable Systems**

# Project Title: Gram-Setu (Unified Rural Engine)

## Vision

Gram-Setu is a **Unified Rural Ecosystem (URE)** platform designed to bridge the digital divide for India's heartland. Instead of isolated tools, we provide a **Multi-Agent "Operating System"** that manages the three most critical pillars of rural life: **Agricultural Health, Government Welfare Access, & Sustainable Resource Management**.

## Core Solution

Our solution employs a **Supervisor-Worker Architecture** to act as a proactive digital companion for rural citizens:

- **Orchestrator:** A central "Supervisor" agent that manages intent and conversation state, ensuring that data flows seamlessly between domains.
- **Agri-Expert :** Provides instant crop disease diagnosis using multimodal vision grounded in the Plant Village dataset and connects users to real-time Mandi prices.
- **Policy-Navigator :** Simplifies complex government schemes (like PM-Kisan and PKVY) using RAG (Retrieval-Augmented Generation) to provide eligibility hallucination-free advice.
- **Eco-Resource :** Processes soil and weather data to provide precise irrigation and energy schedules, significantly reducing resource waste.

## Technical Architecture

- **Intelligence:** Powered by **Claude 3.5 Sonnet** via Amazon Bedrock for high-fidelity reasoning.
- **Orchestration:** Built using the **Strands Agents SDK** for modular, agentic workflows.
- **Data Strategy:** Uses **Amazon S3** for image and sensor log storage, with **Bedrock Knowledge Bases** for grounding policy advice in verified datasets.
- **Responsible AI:** Integrated with **Bedrock Guardrails** to ensure all advice is safe, culturally sensitive, and **privacy-compliant**.
- **MCP:** It allows orchestrator to securely connect to external data source and act as a standardized **plug and play** interface between Agent and data

## Impact & Scalability

Gram-Setu is built to be voice-first and multilingual, making it accessible to users regardless of their literacy level. By optimizing resource use and providing direct market/subsidy linkages, the platform aims to increase net farmer income by **18–22%** while reducing water consumption by up to **45%**.

# Your solution should be able to explain the following

## How different is it from any of the other existing ideas?

Unlike existing siloed agri-tech solutions, our platform unifies multiple rural services into a single AI-driven ecosystem that delivers integrated, actionable decisions

- Potential for AI-driven unified decision support not just for specific tasks but for overarching strategy
- Holistic integration of data across advisory, environmental, market, and operational layers
- A shift from standalone tools or service networks to intelligent ecosystem insights

## How will it be able to solve the problem?

By integrating fragmented rural data and applying AI-driven decision intelligence, the platform transforms complex agricultural challenges into clear solutions.

- Data Integration from Multiple Sources
- AI-Driven Analysis & Intelligence
- Actionable Recommendations
- Trade-off Based Decision Making
- Continuous Learning & Improvement

## USP of the proposed solution

A unified AI decision ecosystem that transforms fragmented rural data into actionable, intelligent decisions

- One Platform, One AI Brain
- Trade-Off Aware Decision Intelligence
- Ecosystem-Scale Impact

# Core Differentiator

 Government of India  
Gram-Sewa  
URE

## Competitive Positioning: Unified Rural Ecosystem (URE)

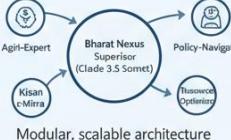
Targeting 150 Million Smallholder Farmers in India

Target Customer Segments	Market Category
 <ul style="list-style-type: none"> <li>Smallholder Farmers (1.08 hectares avg.)</li> <li>Low-Literacy Users (Voice-First)</li> <li>12+ Indian Regional Languages</li> <li>Service-Integrators</li> </ul>	 <ul style="list-style-type: none"> <li>AI-Powered Agri-Advisory</li> <li>Multi-Agent Converstion AI</li> <li>Rural Digital Ecosystems</li> </ul>

Competitor Type Landscape	Competitive Landscape	URE's Advantage
Direct (Agri-Tech)	E-commerce & siloed AI advisory	 Agentic Orchstration: Specialized, collaborative agents
Direct (Dawant (Data-as-Service))	E-commerce & siloed AI advisory	 Proactive Intelligence: IOT triggers & realme alerts
Indirect (Govt. e-Mitra)	Satellite imagry & NDVI analysis	 Contextual Cross-talk: Links as diagnosis to subsidies

### ★ Core Differentiators

1. Supervisor-Worker Agent Orchestration



Modular, scalable architecture

1.  →  Auto-links insights (e.g., Pest → Subsy)

2. Contextual Synthesis (Cross-Talk)

3.  IOT triggers alerts BEFORE crop stress 

4. Amazied RAG & Safety Guardrails Bedrock: Factual Accuracy & Privacy

### Unique Value Proposition (UVP)


URE: A proactive, synthsizez "Digital Village Companion" that anticipates needs and coordinates expert solutions.

 Strategic Moats


- Architectural Complexity
- Services-First Trust

 Voice-First Trust

## List of features offered by the solution



## Process flow diagram



# Wireframes diagrams of the proposed solution

Language selection option

Language: English

URE: Unified Rural Ecosystem

Application name

New Chat

History: Wheat Disease

History: PM-Kisan Status

Settings

Navigation to load  
older communication

Namaste! I am your Unified Rural Ecosystem assistant. You can send me a photo of your crop, ask about government schemes, or get irrigation advice. How can I help you today?

Agent output providing details of the uploaded images or text, action that needs to be taken for crop and market analysis along with the subsidy information

I found these yellow spots on my wheat leaves. What should I do?

(Image Uploaded: wheat\_leaf.jpg)

I've analyzed your request. Here is what our experts found:

Agri-Expert Agent (Vision Analysis)

Diagnosis: Yellow Rust (*Puccinia striiformis*).

Action: Spray Propiconazole 25 EC @ 1ml/liter of water.

Market: Current wheat price in your local Mandi is ₹2,125/quintal.

Policy-Navigator Agent (Scheme Check)

Subsidy Alert: You are eligible for a 50% subsidy on fungicides under the National Food Security Mission (NFSM). Visit your local block office with your Aadhaar card.

Would you like me to help you apply for the subsidy?

Crop image upload

User voice input

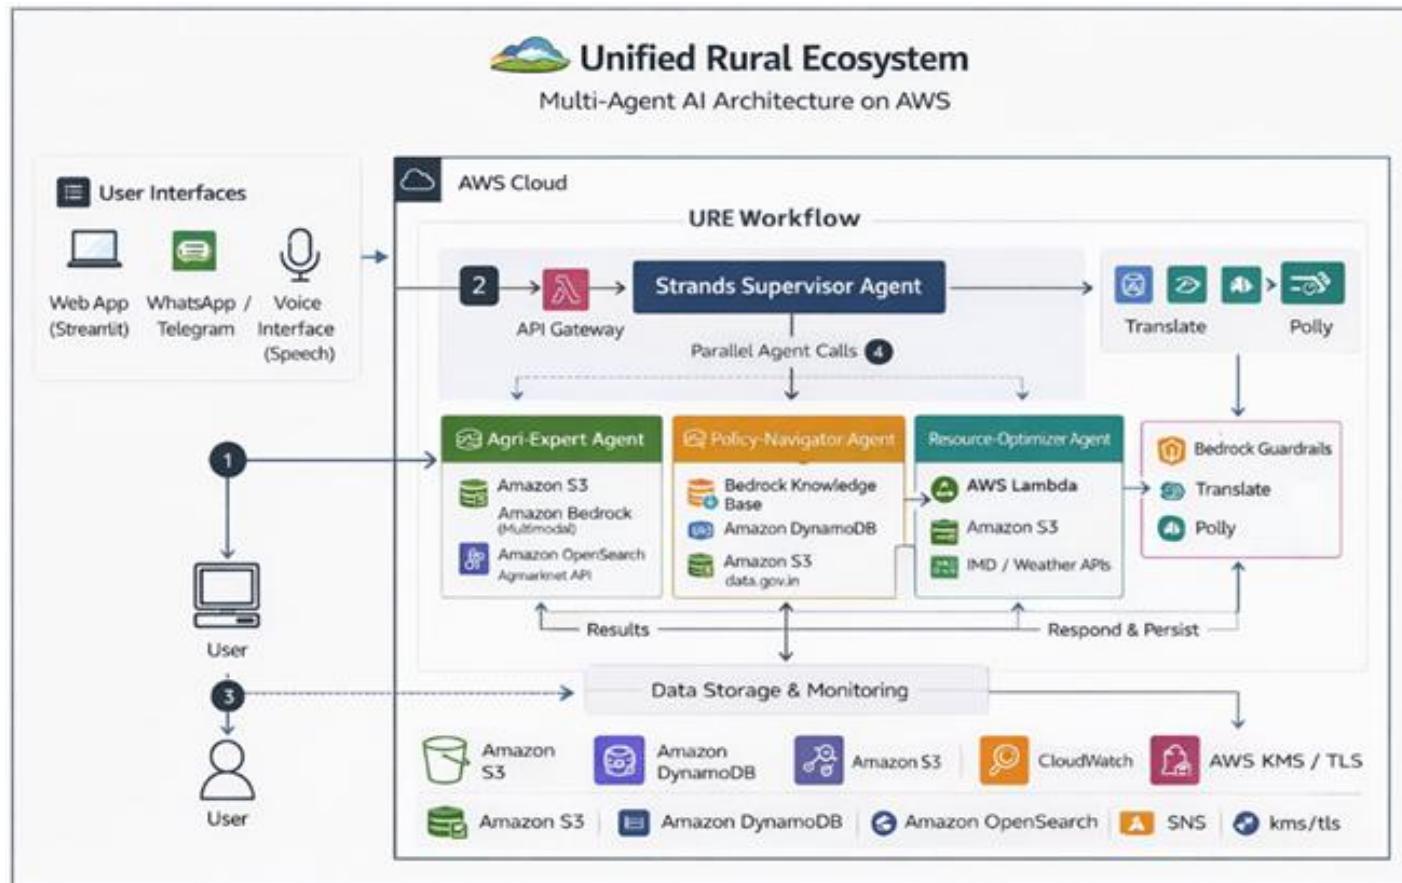
User message



Type your query here...

Send

## Architecture diagram of the proposed solution



## Technologies to be used in the solution

### 1. Core Intelligence Layer

- **Amazon Bedrock:** Hosts and manages foundation models securely at scale
- **Anthropic Claude 3.5 Sonnet:** Handles advanced reasoning and decision-making
- **Amazon Titan Embeddings:** Generates vector embeddings for semantic search

### 2. Orchestration & Agents Layer

- **Strands Agents SDK:** An open-source framework used to build the Supervisor-Worker agentic architecture
- **AWS Lambda:** Provides serverless compute to execute logic and connect agents to external APIs

### 3. Data & Knowledge (RAG) Layer

- **Amazon Bedrock Knowledge Bases:** Manages the RAG pipeline to link agents to private data source
- **Amazon OpenSearch Serverless:** The vector database for low-latency retrieval of government schemes
- **Amazon S3:** Scalable object storage for the Plant Village dataset and user-uploaded images
- **Amazon DynamoDB:** Maintains session state and metadata.

### 4. Interface

- **Streamlit:** User-facing application interface
- **Python Runtime:** Core execution environment for agent workflows
- **Kiro Tool:** Supports agent monitoring and debugging
- **MCP :** Standardized "plug-and-play" interface between Agents



## Estimated implementation cost (1/2)

Layer	Service	Cost Model	Estimated Cost
Core Intelligence	Amazon Bedrock (Claude 3.5 Sonnet)	Pay-per-inference (tokens)	Moderate - High
Core Intelligence	Amazon Titan Embeddings	Pay-per-embedding	Low
Orchestration & Agents	Strands Agents SDK	Open-source	Free
Orchestration & Agents	AWS Lambda	Pay-per-invocation	Very Low
Data & Knowledge (RAG)	Bedrock Knowledge Bases	Usage-based	Low
Data & Knowledge (RAG)	Amazon Open Search Serverless	Storage + query-based	Moderate
Data & Knowledge (RAG)	Amazon S3	Pay-per-storage	Very Low
Data & Knowledge (RAG)	Amazon DynamoDB	On-demand reads/writes	Low
Interface	Streamlit	Open-source	Free
Interface	Python Runtime	Open-source	Free
Interface	Kiro Tool	Tooling / Dev support	Minimal
Model Context Protocol	Secure & Seamless connection	Open-source	Free

### Notes :

1. Total cost is directly proportional to the number of users.
2. Depending on LLM selected cost can be further optimized

## Estimated implementation cost (2/2)

### Cost Analysis & Optimization

Live financial simulator for the Gram-Setu Multi-Agent Ecosystem (Hackathon Preview)

Scale (Active Users)

600

#### OPTIMIZATION STRATEGY

##### Prompt Caching

Up to 90% discount on repeated context—critical for conversation history.

##### Batch Mode

50% discount for non-real-time tasks (e.g., daily report generation).

##### Model Tiering

Use Claude 3.5 Haiku for simpler queries, reserving Sonnet for complex tasks.

##### Routing Layer

Pre-classify queries to bypass Supervisor for FAQ-style questions, reducing agentic overhead.

#### STRATEGY IMPACT ANALYSIS


- Efficiency Gain: 0.0%
- Opex Avoidance: \$0/mo
- Scale Readiness: Growth Phase

EST. MONTHLY OPEX

**\$690**

UNIT COST / USER

**\$1.15**



### Cost Analysis & Optimization

Live financial simulator for the Gram-Setu Multi-Agent Ecosystem (Hackathon Preview)

Scale (Active Users)

600

#### OPTIMIZATION STRATEGY

##### Prompt Caching

Up to 90% discount on repeated context—critical for conversation history.

##### Batch Mode

50% discount for non-real-time tasks (e.g., daily report generation).

##### Model Tiering

Use Claude 3.5 Haiku for simpler queries, reserving Sonnet for complex tasks.

##### Routing Layer

Pre-classify queries to bypass Supervisor for FAQ-style questions, reducing agentic overhead.

#### STRATEGY IMPACT ANALYSIS

- Efficiency Gain: 38.0%
- Opex Avoidance: \$262.2/mo
- Scale Readiness: Growth Phase

EST. MONTHLY OPEX

**\$427.8**

UNIT COST / USER

**\$0.71**



#### Strategy Definitions

- Prompt Caching: Reuses cached token IDs to bypass compute.
- Batch Mode: Leverages asynchronous processing for 50% API cost reduction.
- Model Tiering: Cost-optimal routing between Claude model variants.
- Routing Layer: Heuristic-based shortcutting to minimize LLM inference.

Please click here to view the  
Cost Analysis

\*\*it's only for simulation purposes actual cost may differ on case to case basis



Scan QR  
Code to view  
on Mobile

Innovation partner **H2S**

Media partner **YOURSTORY**

# AI for Bharat Hackathon

Powered by 

Thank You

